

Technical Report

Linear Regression Models To Predict Cancer Death Rate

CS452 Homework1

Ordinary Least Squares / Gradient Descent / KNN-Regression

ALKIN KORKUT
REYHAN AYDOĞAN
26 November 2023

CONTENTS

1. *Introduction*
 - *Purpose*
 - *Loading and Reading Dataset*
2. *Data Processing and Visualizing*
 - *Preprocessing*
 - *Dataset Description*
 - *Data Visualization*
 - *Dataset Splitting*
 - *Removing Outliers*
 - *Scaling*
3. *Ordinary Least Squares*
 - *Implementation*
 - *Model Evaluation*
 - *Results and Explanation*
4. *Gradient Descent*
 - *Implementation*
 - *Model Evaluation*
 - *Results and Explanation*
5. *KNN - Regression*
 - *Implementation*
 - *Model Evaluation*
 - *Results and Explanation*

1. Introduction

Purpose

The primary goal of this work is to create linear regression models to predict cancer death rate from "cancer.csv" file by using ordinary least squares, gradient descent, and knn - regression approaches. In ordinary least squares method, LinearRegression method is used from Scikit-learn library. As for gradient descent approach, it is written a gradient descent algorithm from scratch. Lastly, knn-regression method is used again from Scikit-learn library. Only NumPy, Pandas, Scikit-learn, Matplotlib, Seaborn, and Chardet libraries are used.

Loading and Reading Dataset

First of all, I have encountered an error while trying to read the 'cancer.csv' file. The error is:

"UnicodeDecodeError: 'utf-8' codec can't decode byte 0xf1 in position 41137: invalid continuation byte"

It is about an issue decoding the file as UTF-8. I tried to solve this problem importing 'chardet' library and then tried to open the dataset. 'Chardet' library already installed in the python. Therefore, I did not need to install again.

```
import chardet
```

```
# Detect encoding
with open("cancer.csv", 'rb') as f:
    result = chardet.detect(f.read())

# Load the dataset with detected encoding
df = pd.read_csv("cancer.csv", encoding=result['encoding'])
```

2. Data Processing and Visualizing

Preprocessing

Missing data could affect the performance of the models. So, it needs to be handled. Also, we need to remove non-numeric features of the data such as categorical or text data. At the first step, I removed the non-numeric features (`#binnedinc`, `#geography`). After purify the data from non-numeric features, I removed the rows with missing data. While performing these operations, I checked the shape of the data frame in the intermediate steps and checked whether there was null data.

Dataset Description

Before the preprocessing, we have 3047 rows and 34 columns at the data frame. Afterwards, we have only 591 rows and 32 columns. Following statistics are available when the code is run.

- I. Number of samples.
- II. Number of features.
- III. Mean and Variance for each feature.
- IV. Correlation matrix.

These stats can provide valuable and useful insights for the characteristics of the data. Number of samples includes the total number of instances after preprocessing. It gives an understanding of the dataset's size. Number of features indicates the attributes of the data. It gives an understanding about the dimensionality of the dataset. High variance may indicate significant variability in data points for that feature. When there is low variance, it means that the data points are near the mean. The correlation matrix illustrates the degree of relationship between two features. A high positive correlation between features indicates a tendency for them to rise or fall together. On the other hand, an inverse relationship is suggested by a high negative correlation. Additionally, a low correlation or close to zero indicates a weak or no linear relationship.

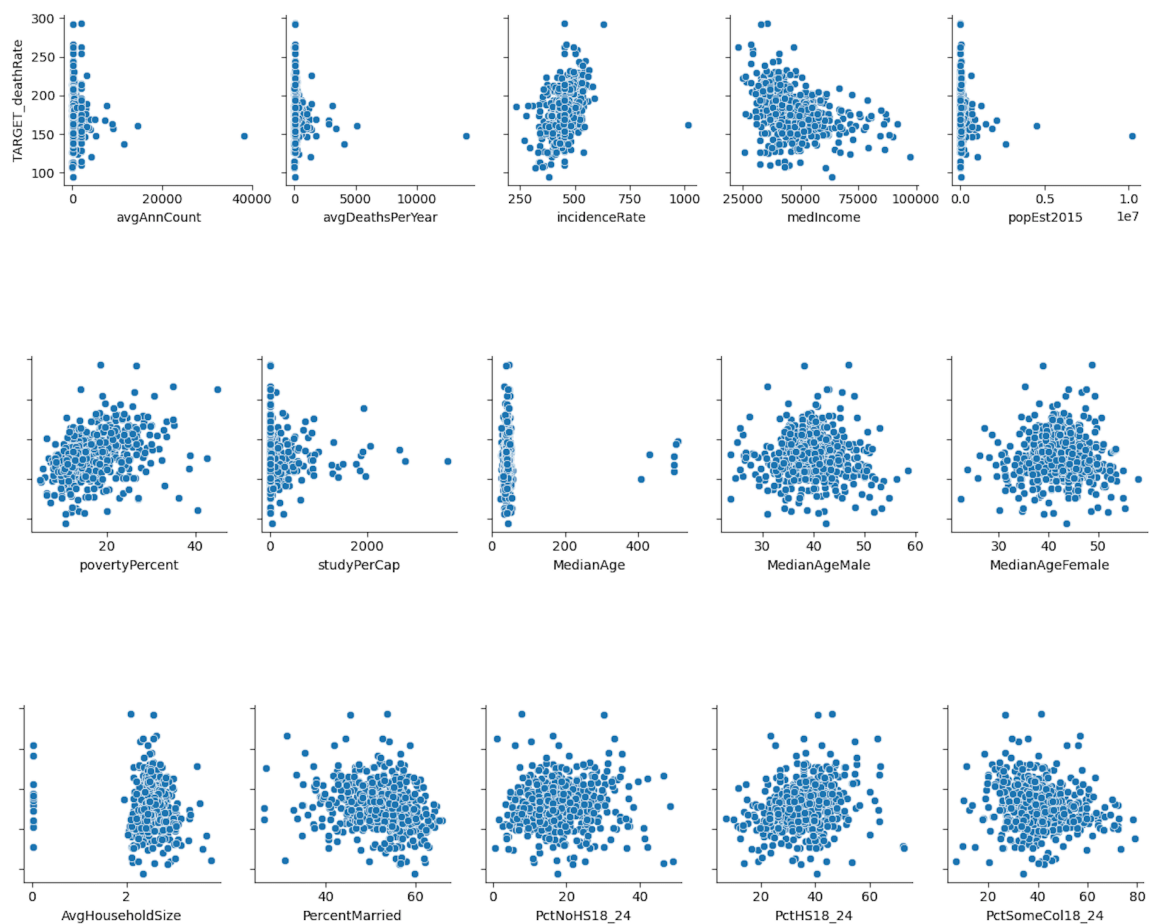
Data Visualization

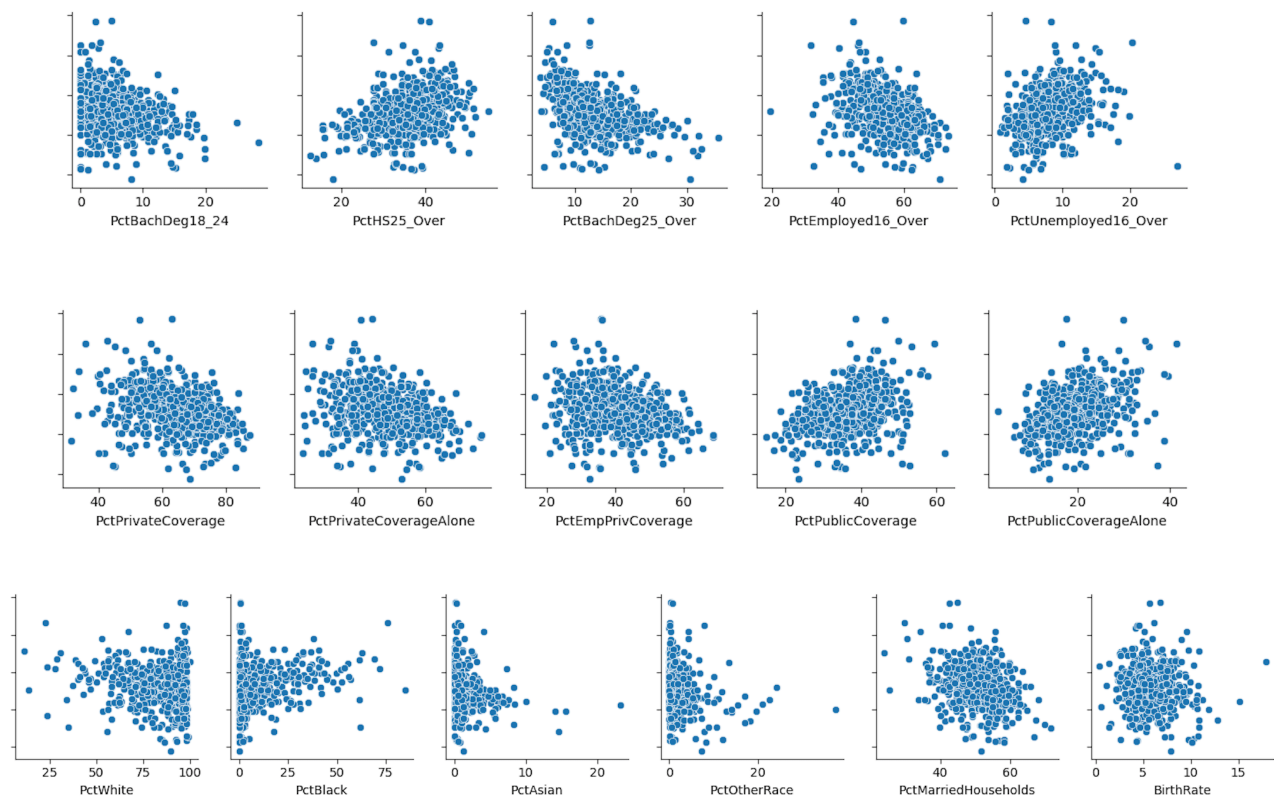
Before start plotting, first I moved the y_target_variable (TARGET_deathRate) column to the last index of the data frame. After this moving operation I created scatter diagrams to visualize the relationship between the death rate (the target) and other features separately.

```
targetColumn = 'TARGET_deathRate'
# Move the target column to the last index of the dataframe
df = df[[col for col in df.columns if col != targetColumn] + [targetColumn]]

# Create scatter diagrams to visualize the relationship between
# the death rate (the target) and other features separately
sns.pairplot(df, x_vars=df.columns[:-1], y_vars='TARGET_deathRate', kind='scatter')
plt.show()
```

Scatter Diagrams:





Dataset Splitting

Dataset is splitted as test_size = 0.20 and random_state = 271.

Removing Outliers

I removed the outliers using Z-Score technique.

```
# Removing outliers using Z-score technique
z_scores = (X_train - np.mean(X_train, axis=0)) / np.std(X_train, axis=0)
outliers = (np.abs(z_scores) > 3).all(axis=1)
X_train_no_outliers = X_train[~outliers]
y_train_no_outliers = y_train[~outliers]
```

Using this procedure I tried to provide the dataset free from the impact of extreme values. It's a widely used method to ensure that statistical measures fairly represent the vast majority of the data and to clean up data.

Scaling

```
from sklearn.preprocessing import StandardScaler  
  
# Scaling using StandardScaler  
scaler = StandardScaler()  
X_trainScaled = scaler.fit_transform(X_train_no_outliers)  
X_testScaled = scaler.transform(X_test)
```

The aim here is to bring the features onto the same scale. The performance and convergence speed of machine learning algorithms can be enhanced by features that are nearly normally distributed and/or on similar scales. By using standardization technique the feature columns are centered at mean 0 with standard deviation 1 so that the feature columns take the form of a normal distribution.

3. Ordinary Least Squares

Implementation

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

def ordinary_least_squares(X_train, X_test, y_train, y_test, numFeatures):
    # Select the first 'num_features' features.
    selectedFeatures = X_train[:, :numFeatures]

    # Fit the model with LinearRegression().
    model = LinearRegression()
    model.fit(selectedFeatures, y_train)

    # Make predictions on the test set.
    selectedFeaturesTest = X_test[:, :numFeatures]
    y_pred = model.predict(selectedFeaturesTest)

    # Evaluate the model by calculating mse, mae, and r2 score.
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    return model, mse, mae, r2

results = {}

# Evaluating models for F = 5, 10, 20 values
for numOffFeatures in [5, 10, 20]:
    model, mse, mae, r2 = ordinary_least_squares(X_trainScaled, X_testScaled, y_train, y_test, numOffFeatures)

    # Storing results in the dictionary
    results[numOffFeatures] = {
        'f_value': numOffFeatures,
        'model': model,
        'mse': mse,
        'mae': mae,
        'r2': r2
    }

    print(f"\nResults for {numOffFeatures} features:")
    print(f"MSE: {mse}")
    print(f"MAE: {mae}")
    print(f"R2 Score: {r2}")
    print(f"Coefficients: {model.coef_}")
    print(f"Intercept: {model.intercept_}")

# Identifying the best model (The best model is with lowest MSE)
bestModelKey = min(results, key=lambda k: results[k]['mse'])
bestModel = results[bestModelKey]['model']

print("\n")
print("---- RESULT ----")
# Printing the best model
print("The Best Model with f value", results[bestModelKey]['f_value'])
print(f"{bestModelKey} with mse:", results[bestModelKey]['mse'])
print("with mae:", results[bestModelKey]['mae'])
print("with R2 score:", results[bestModelKey]['r2'])

# Printing the linear equation for the best model
print("\nLinear Equation for the Best Model:")
equation = f"y = {bestModel.intercept_} "
for i, coef in enumerate(bestModel.coef_):
    equation += f"+ ({coef} * X_{i}) "
print(equation)
```


Model Evaluation

Results for 5 features:

MSE: 463.51974876777234

MAE: 16.07056016183814

R2 Score: 0.3100302483143771

Coefficients: [-4.79333086 10.0087527 10.12404938 -11.27436885
-5.35973882]

Intercept: 179.33326271186442

Results for 10 features:

MSE: 443.1639611822952

MAE: 15.64879809384971

R2 Score: 0.34033074304637934

Coefficients: [-4.37843822 8.99460824 10.25193451 -7.60562729
-5.0248862 4.01793302

-0.44724573 -0.82089988 -2.19776472 3.79246302]

Intercept: 179.33326271186442

Results for 20 features:

MSE: 446.0138685405535

MAE: 15.308115424103127

R2 Score: 0.3360885292517528

Coefficients: [-1.79565472e+00 7.15418346e+00 9.31413288e+00
-9.13674545e-01
-4.42582484e+00 2.14458324e+00 3.57875570e-02 -1.22942575e+00
-4.95286280e-01 -3.42388629e-02 -3.28715947e-01 -7.88766000e-01
-7.28401451e+01 -8.03662482e+01 -9.88841428e+01 -4.05484774e+01
4.49185785e+00 -4.33834617e+00 -1.39613724e+00 3.20616655e+00]

Intercept: 179.33326271186442

MSE measures the average squared difference between actual and predict values. When we look at the MSE results, lower mean-squared-error indicates better model performance. But we should know it MSE is sensitive to outliers since it squares the differences. On the other hand mean-absolute-

error is less sensitive to outliers comparing to MSE. MAE measures the average absolute difference between actual and predicted values. By looking at MAEs we should choose with the lowest MAE if we want a metric that gives equal weight to all errors. Other than MSE and MAE metrics we have R2 score. It represents the proportion of the variance in the dependent variable that is predictable from the independent variables. It ranges from 0 to 1, where 1 indicates a perfect fit. In the case of choosing best model with R2 score, we should consider the model with the highest R2 score. But at the same time, we should be careful with overfitting issue. You can see the best method that I chose under the next subtitle after applying ordinary-least-squared method.

Results and Explanation

The Best Model with f value 10:

LinearRegression() with

Mean-Squared-Error: 443.1639611822952

Mean-Absolute-Error: 15.64879809384971

R2 Score: 0.34033074304637934

Linear Equation for the Best Model:

$$y = 179.33326271186442 + (-4.378438224533927 * X_0) + (8.994608244325995 * X_1) + (10.251934509971468 * X_2) + (-7.6056272879295195 * X_3) + (-5.024886200394534 * X_4) + (4.017933015095547 * X_5) + (-0.4472457286448114 * X_6) + (-0.8208998761859805 * X_7) + (-2.197764718164857 * X_8) + (3.792463024046885 * X_9)$$

When I look at the results of the metrics of the three models, the model with f value 10 has lowest MSE and highest R2 score. Since we want lower MSE, MAE and higher R2 score it can be a good choice among three models. When look at the MAE for the model with f value 10, it is not the lowest among three model but actually it is not a bad value when comparing all of them. Also we have better results with MSE and R2 score other than the other models. So the best model with f value 10 is the best model that I chose.

4. Gradient Descent

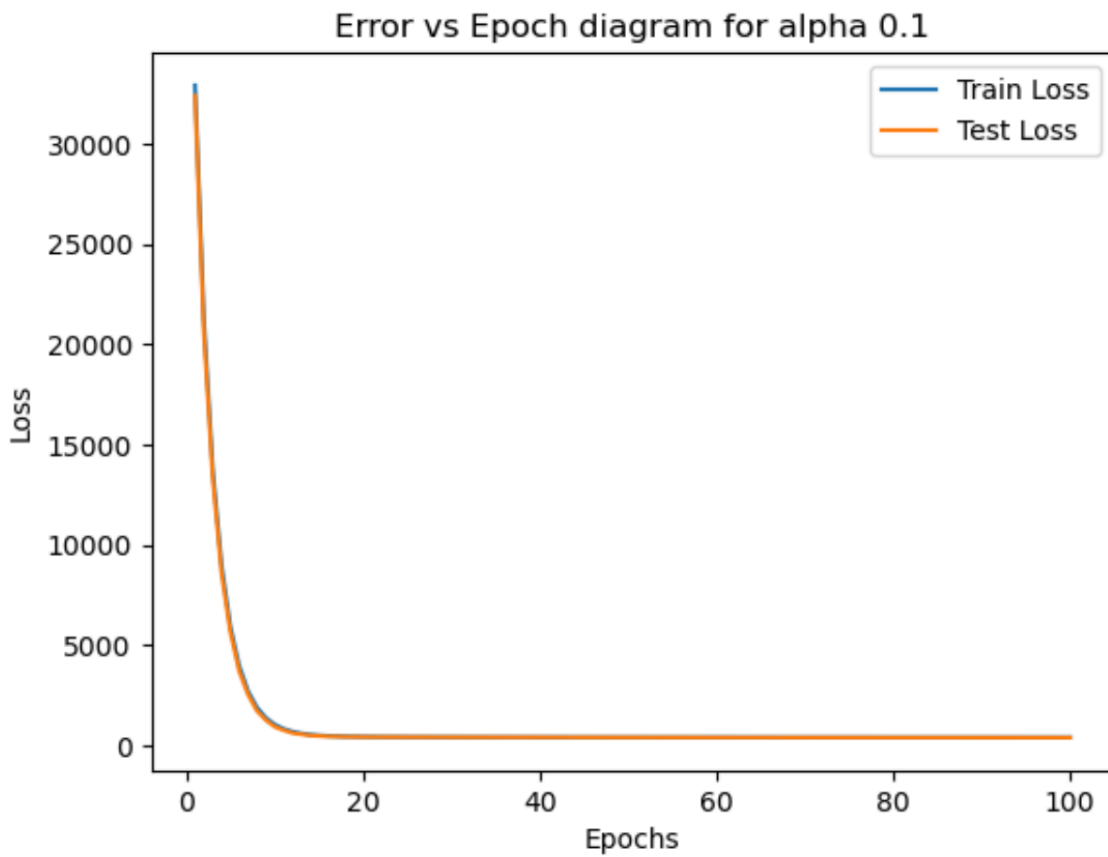
Implementation

```
def gradient_descent(X_train, X_test, y_train, y_test, alpha, numOfIterations):  
    # 'Weights' correspond to coefficients for the features in the input data.  
    # 'Bias' corresponds to intercept in the linear equation  
    weights = np.zeros(X_train.shape[1])  
    bias = 0  
  
    # 'Epochs' correspond to the number of times the gradient descent  
    # algorithm will iterate through the entire dataset  
    epochs = numOfIterations  
  
    # Declaring numpy arrays in order to store the error  
    # losses calculated for training and test sets here  
    trainLosses = []  
    testLosses = []  
  
    # Each time iterate through the entire dataset  
    for ep in range(epochs):  
        # Calculating predictions and loss for the training set  
        yTrainPred = np.dot(X_train, weights) + bias  
        trainLoss = np.mean(np.square(yTrainPred - y_train))  
        trainLosses = np.append(trainLosses, trainLoss)  
  
        # Calculating predictions and loss for the test set  
        yTestPred = np.dot(X_test, weights) + bias  
        testLoss = np.mean(np.square(yTestPred - y_test))  
        testLosses = np.append(testLosses, testLoss)  
  
        # Updating weights and bias using gradient descent in  
        # the opposite direction of the gradients in order  
        # to minimize the loss  
        weights = weights - alpha * (2/len(X_train)) * np.dot(X_train.T, (yTrainPred - y_train))  
        bias = bias - alpha * (2/len(X_train)) * np.sum(yTrainPred - y_train)  
  
    return weights, bias, trainLosses, testLosses  
  
# Train models with different learning rates  
learningRates = [0.1, 0.5, 0.01]  
  
for alpha in learningRates:  
    weights, bias, trainLosses, testLosses = gradient_descent(X_trainScaled, X_testScaled, y_train, y_test,  
                                                             alpha, 100)  
  
    print(f"Results for learning rate {alpha}:")  
    print(f"Weights: {weights}, Bias: {bias}")  
    # Plot the epoch-loss diagram  
    plt.plot(range(1, len(trainLosses)+1), trainLosses, label='Train Loss')  
    plt.plot(range(1, len(testLosses)+1), testLosses, label='Test Loss')  
    plt.xlabel('Epochs')  
    plt.ylabel('Loss')  
    plt.title(f'Error vs Epoch diagram for alpha {alpha}')  
    plt.legend()  
    plt.show()  
    print(f"----Losses for train data for the model with alpha val: {alpha}----")  
    print(trainLosses)  
    print("\n")  
    print(f"----Losses for test data for the model with alpha val: {alpha}----")  
    print(testLosses)  
    print("-----\n")
```

Model Evaluation

Results for learning rate 0.1:

Weights: [-0.92531662 3.09442635 8.35931988 0.64072013 -1.27041053 -0.30584311
-0.32105436 -1.24775457 -0.85493468 -0.97927522 0.53846628 6.64534416
-1.97006177 2.55388191 0.32916207 -3.00216434 3.19750232 -5.20617308
-4.20315507 2.67502082 -4.54455691 -1.58242223 5.68186561 -2.65197347
2.0254707 3.0561517 2.57411658 -0.06229792 -3.26228576 -8.85048084
-0.22830118], Bias: 179.3332626753336



----Losses for train data for the model with alpha val: 0.1----


[32920.56489407	21234.79277401	13764.27188872	8984.01177293
5923.16770606	3962.01516558	2704.53932501	1897.52206472
1378.98575924	1045.2850814	830.08166426	690.90336955
600.54913619	541.59119701	502.85785034	477.18295149
459.96570787	448.24888533	440.12871977	434.37709433
430.1995659	427.08034557	424.6829447	422.78646287
421.24470207	419.95990372	418.86585685	417.91701499
417.08146817	416.33639063	415.66508025	415.05502294
414.49661852	413.98233458	413.50613833	413.06310996
412.64917518	412.26091659	411.89543802	411.55026446
411.22326681	410.91260368	410.61667571	410.33408879
410.06362411	409.8042133	409.55491768	409.31491068
409.08346298	408.85992963	408.6437392	408.4343842
408.23141303	408.0344228	407.84305327	407.65698149
407.47591719	407.29959877	407.12778975	406.96027579
406.79686199	406.6373706	406.48163898	406.32951786
406.18086982	406.03556789	405.89349444	405.75454012
405.61860297	405.48558764	405.35540469	405.22796999
405.10320421	404.98103231	404.86138317	404.74418922
404.6293861	404.51691239	404.40670939	404.29872086
404.19289282	404.08917343	403.98751277	403.88786274
403.79017694	403.69441052	403.60052014	403.50846383
403.41820092	403.32969199	403.24289877	403.15778412
403.07431192	402.99244709	402.91215547	402.83340384
402.75615984	402.68039194	402.60606944	402.53316238]

----Losses for test data for the model with alpha val: 0.1----

[32401.88453782	20656.25079738	13391.26263086	8598.52586271
5670.13409717	3704.84933614	2539.58048971	1730.11449335
1274.7240024	938.24832606	765.54209492	623.19864707
561.24500207	499.01887152	479.29752639	450.47162791
446.083954	431.4477072	432.09609297	423.69971833
425.61400161	420.12875747	422.04504359	418.04620976
419.62352585	416.4768153	417.68769028	415.09166568
415.99669098	413.79385433	414.46435529	412.56244585
413.05815757	411.39678161	411.76254507	410.29855347
410.56662149	409.26732649	409.460609	408.30035105
408.43519721	407.39340327	407.48166319	406.54161702
406.59207515	405.74006367	405.75940512	404.98408841
404.97754117	404.26947115	404.24123123	403.59247722
403.54599103	402.94984589	402.8879995	402.33874814
402.26399596	401.7567326	401.67118656	401.20167008
401.10716293	400.67170219	400.56983368	400.16519619
400.05736783	399.68070678	399.56814876	399.21694429
399.10073704	398.77274872	398.65384067	398.34706871
398.22629117	397.93894454	397.81702454	397.54749446
397.42506597	397.17190371	397.04951751	396.81141562
396.68954811	396.46532441	396.3443855	396.13296932
396.01330956	395.81372975	395.69564672	395.50702125
395.39076539	395.21229218	395.09807201	394.92902083
394.8170077	394.656713	394.54704534	394.39489988
394.28768712	394.14313623	394.03846223	393.90099874]

Results for learning rate 0.5:

Weights: [2.60644398e+93 2.19856606e+93 6.13395671e+92 6.08415788e+93
2.16445955e+93 -5.79491680e+93 1.19916945e+93 -3.26268814e+92
-9.49942622e+92 -1.46442328e+93 2.79880262e+92 2.46132430e+93
-3.59961421e+93 -3.01519696e+93 3.37001344e+93 4.51302734e+93
-3.45396704e+93 5.52850764e+93 5.91273898e+93 -4.32691273e+93
6.34976602e+93 6.63613473e+93 6.10154142e+93 -6.18290553e+93
-6.20648346e+93 2.45284622e+93 -2.68093589e+93 3.22018795e+93
1.01270930e+92 3.05933231e+93 -6.33129499e+92], Bias: 2.6682362639989975e+77



```

----Losses for train data for the model with alpha val: 0.5----
[3.29205649e+004 1.16429103e+004 8.36009003e+005 6.26290834e+007
 4.69552586e+009 3.52060538e+011 2.63970445e+013 1.97922120e+015
 1.48399898e+017 1.11268677e+019 8.34280819e+020 6.25534973e+022
 4.69019537e+024 3.51665910e+026 2.63675397e+028 1.97701037e+030
 1.48234156e+032 1.11144409e+034 8.33349079e+035 6.24836366e+037
 4.68495729e+039 3.51273165e+041 2.63380921e+043 1.97480241e+045
 1.48068606e+047 1.11020282e+049 8.32418383e+050 6.24138539e+052
 4.67972506e+054 3.50880858e+056 2.63086773e+058 1.97259692e+060
 1.47903240e+062 1.10896292e+064 8.31488725e+065 6.23441492e+067
 4.67449867e+069 3.50488989e+071 2.62792954e+073 1.97039390e+075
 1.47738060e+077 1.10772442e+079 8.30560106e+080 6.22745223e+082
 4.66927812e+084 3.50097557e+086 2.62499462e+088 1.96819333e+090
 1.47573064e+092 1.10648730e+094 8.29632525e+095 6.22049732e+097
 4.66406340e+099 3.49706563e+101 2.62206299e+103 1.96599522e+105
 1.47408252e+107 1.10525155e+109 8.28705979e+110 6.21355018e+112
 4.65885450e+114 3.49316006e+116 2.61913463e+118 1.96379957e+120
 1.47243624e+122 1.10401719e+124 8.27780467e+125 6.20661079e+127
 4.65365142e+129 3.48925884e+131 2.61620954e+133 1.96160637e+135
 1.47079180e+137 1.10278421e+139 8.26855990e+140 6.19967916e+142
 4.64845416e+144 3.48536199e+146 2.61328772e+148 1.95941562e+150
 1.46914920e+152 1.10155260e+154 8.25932545e+155 6.19275526e+157
 4.64326269e+159 3.48146948e+161 2.61036916e+163 1.95722732e+165
 1.46750843e+167 1.10032237e+169 8.25010131e+170 6.18583910e+172
 4.63807703e+174 3.47758132e+176 2.60745386e+178 1.95504145e+180
 1.46586950e+182 1.09909352e+184 8.24088747e+185 6.17893066e+187]

```

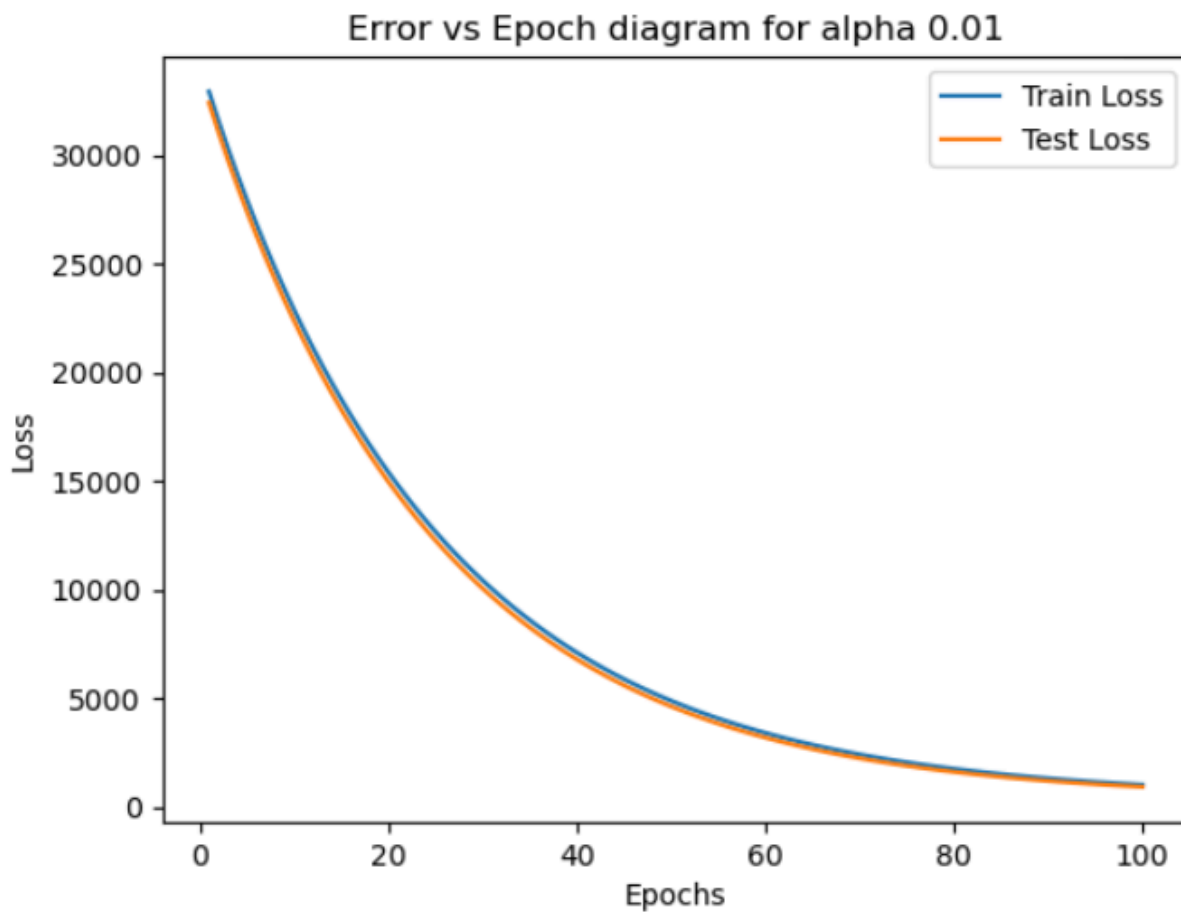
```

----Losses for test data for the model with alpha val: 0.5----
[3.24018845e+004 1.13419999e+004 8.21452994e+005 6.13834578e+007
 4.60811077e+009 3.45843909e+011 2.59448760e+013 1.94583635e+015
 1.45915438e+017 1.09412526e+019 8.20387303e+020 6.15126174e+022
 4.61218113e+024 3.45817543e+026 2.59290728e+028 1.94413591e+030
 1.45769311e+032 1.09296311e+034 8.19492294e+035 6.14446715e+037
 4.60705685e+039 3.45432275e+041 2.59001484e+043 1.94196586e+045
 1.45606556e+047 1.09174262e+049 8.18577126e+050 6.13760511e+052
 4.60191169e+054 3.45046493e+056 2.58712228e+058 1.93979705e+060
 1.45443940e+062 1.09052335e+064 8.17662927e+065 6.13075054e+067
 4.59677221e+069 3.44661140e+071 2.58423294e+073 1.93763065e+075
 1.45281506e+077 1.08930544e+079 8.16749749e+080 6.12390362e+082
 4.59163846e+084 3.44276218e+086 2.58134683e+088 1.93546668e+090
 1.45119254e+092 1.08808889e+094 8.15837590e+095 6.11706436e+097
 4.58651045e+099 3.43891725e+101 2.57846395e+103 1.93330512e+105
 1.44957182e+107 1.08687369e+109 8.14926451e+110 6.11023273e+112
 4.58138817e+114 3.43507661e+116 2.57558428e+118 1.93114598e+120
 1.44795292e+122 1.08565985e+124 8.14016329e+125 6.10340873e+127
 4.57627161e+129 3.43124027e+131 2.57270783e+133 1.92898924e+135
 1.44633582e+137 1.08444737e+139 8.13107223e+140 6.09659235e+142
 4.57116076e+144 3.42740821e+146 2.56983459e+148 1.92683492e+150
 1.44472054e+152 1.08323625e+154 8.12199133e+155 6.08978359e+157
 4.56605562e+159 3.42358043e+161 2.56696456e+163 1.92468300e+165
 1.44310705e+167 1.08202647e+169 8.11292057e+170 6.08298242e+172
 4.56095618e+174 3.41975692e+176 2.56409774e+178 1.92253349e+180
 1.44149537e+182 1.08081805e+184 8.10385994e+185 6.07618886e+187]

```

Results for learning rate 0.01:

Weights: [0.05267679 0.96669533 7.58546084 -0.85728114 0.11199243 0.71898088
-0.0140828 -0.82632621 -0.74459755 -0.76179112 -0.18090837 -0.26527586
-1.64260107 2.75569092 -0.21300241 -2.44603547 3.79154552 -3.54573718
-1.64426964 2.48563643 -1.00910168 -0.14862673 1.87535418 0.46335149
1.5820054 0.95509913 1.11911036 -0.54268668 -3.08032326 -2.20577531
-0.61503452], Bias: 155.55016505385987




```

----Losses for train data for the model with alpha val: 0.01----
[32920.56489407 31584.80725505 30318.22586744 29112.45910501
 27961.43378387 26860.591853 25806.38593359 24795.94991882
 23826.88357843 22897.11142012 22004.78993107 21148.24635198
 20325.93801488 19536.42510189 18778.35217384 18050.43543921
 17351.45378966 16680.24231591 16035.6874655 15416.72329503
 14822.32845964 14251.52370592 13703.36971511 13176.96519571
 12671.44515922 12185.97933456 11719.77069185 11272.05405523
 10842.09479106 10429.1875617 10032.65513781 9651.84726408
 9286.13957414 8934.9325516 8597.6505345 8273.74076074
 7962.67245269 7663.93593901 7377.04181202 7101.52011913
 6836.91958686 6582.80687606 6338.76586709 6104.39697359
 5879.31648391 5663.15592875 5455.56147431 5256.19333957
 5064.72523696 4880.84383533 4704.24824449 4534.64952023
 4371.77018914 4215.34379242 4065.11444792 3920.83642956
 3782.27376362 3649.19984108 3521.39704545 3398.65639535
 3280.77720147 3167.56673704 3058.83992155 2954.419017
 2854.13333626 2757.81896303 2665.31848299 2576.48072554
 2491.16051597 2409.21843736 2330.52060197 2254.93843183
 2182.34844788 2112.63206767 2045.67541092 1981.36911303
 1919.60814574 1860.29164519 1803.32274656 1748.60842548
 1696.05934562 1645.58971236 1597.11713233 1550.56247843
 1505.84976031 1462.9059999 1421.66111197 1382.04778936
 1344.00139281 1307.45984516 1272.36352973 1238.65519274
 1206.2798496 1175.18469496 1145.31901627 1116.6341108
 1089.08320594 1062.62138269 1037.20550217 1012.79413501]

```

```

----Losses for test data for the model with alpha val: 0.01----
[32401.88453782 31057.81388555 29786.37996763 28578.53506528
 27427.64369862 26328.68984072 25277.75492894 24271.67319667
 23307.80326954 22383.87611765 21497.89325328 20648.05807281
 19832.7291268 19050.38795055 18299.61660552 17579.08173247
 16887.5229994 16223.7445394 15586.60844173 14975.02966901
 14387.97197808 13824.44455819 13283.49919092 12764.22779677
 12265.76027478 11787.26256885 11327.93491349 10887.01022517
 10463.75261371 10057.45599546 9667.44279363 9293.06271506
 8933.69159461 8588.73030039 8257.60369408 7939.75964174
 7634.66807119 7341.82007264 7060.72703963 6790.91984784
 6531.94806949 6283.37922126 6044.79804415 5815.80581336
 5596.01967692 5385.07202159 5182.60986472 4988.29427098
 4801.79979272 4622.81393305 4451.03663051 4286.17976448
 4127.96668053 3976.13173466 3830.41985586 3690.58612605
 3556.39537688 3427.62180249 3304.04858774 3185.46755121
 3071.67880238 2962.49041243 2857.71809812 2757.18491822
 2660.72098192 2568.16316891 2479.35486046 2394.14568123
 2312.39125129 2233.95294793 2158.69767698 2086.49765315
 2017.23018909 1950.77749277 1887.02647297 1825.8685524
 1767.19948825 1710.9191999 1656.93160338 1605.14445244
 1555.46918594 1507.82078122 1462.11761337 1418.28131999
 1376.23667142 1335.91144601 1297.23631042 1260.14470463
 1224.57273149 1190.45905074 1157.74477711 1126.37338262
 1096.29060258 1067.44434549 1039.78460639 1013.26338375
 987.83459957 963.45402275 940.07919541 917.66936219]

```

Results and Explanation

When we look at the diagram and loss values of the model with an alpha value of 0.1, we see that the loss value decreases very quickly in the first iterations. At the same time, by comparing the minimum loss values achieved by all 3 models, the model with the alpha value of 0.1 provides us the lowest loss values at the end of the epochs. Therefore for the best model I chose the model with alpha value of 0.1.

5. KNN - Regression

Implementation

```
from sklearn.neighbors import KNeighborsRegressor

def knn_regression(X_train, X_test, y_train, y_test, k):
    # Applying KNN Regressor
    model = KNeighborsRegressor(n_neighbors=k)
    # Fit the model
    model.fit(X_train, y_train)
    # Predict
    yPred = model.predict(X_test)
    # Calculate MSE
    mse = mean_squared_error(y_test, yPred)

    return model, mse

# K values
kValues = [3, 5, 10]
results = {}
for k in kValues:
    model, mse = knn_regression(X_trainScaled, X_testScaled, y_train, y_test, k)
    # Using dictionary for results
    results[k] = {
        'model': model,
        'mse': mse
    }
    print(f"Results for K={k}:")
    print(f"MSE: {mse}")
    print("=", round(mse))
    print("----")

# Identifying the best model (The best model is with lowest MSE)
bestModelKey = min(results, key=lambda k: results[k]['mse'])
bestModel = results[bestModelKey]['model']

print("---- RESULT ----")
# Printing the best model with k value
print("The Best Model with k value:")
print(f"{bestModel} with mse:", results[bestModelKey]['mse'])
```

Model Evaluation

```
Results for K=3:
MSE: 495.06725490196055
    = 495
----
Results for K=5:
MSE: 478.1488806722688
    = 478
----
Results for K=10:
MSE: 476.60784285714266
    = 477
----
```

Results and Explanation

```
---- RESULT ----
The Best Model with k value:
KNeighborsRegressor(n_neighbors=10) with mse: 476.60784285714266
```

As a result, when we compare all models, we should choose the model with the lowest mean-squared-error value as the best model. Because, lower mean-squared-error indicates better model performance. For this reason, I chose the best model among these 3 models as the model with k value 10 that has the lowest MSE value.